

UNIT – IV

Introduction to Matplotlib

Matplotlib is the most popular plotting library for python which gives control over every aspect of a figure. It was designed to give the end user a similar feeling like MATLAB's graphical plotting. In the coming sections we will learn about Seaborn that is built over matplotlib. The official page of Matplotlib is <https://matplotlib.org>. You can use this page for official installation instructions and various documentation links. One of the most important section on this page is the gallery section - <https://matplotlib.org/gallery.html> - it shows all the kind of plots/figures that matplotlib is capable of creating for you. You can select anyone of those, and it takes you the example page having the figure and very well documented code. Another important page is https://matplotlib.org/api/pyplot_summary.html- and it has the documentation functions in it.

Matplotlib's architecture is composed of three main layers: the back-end layer, the artist layer where much of the heavy lifting happens and is usually the appropriate programming paradigm when writing a web application server, or a UI application, or perhaps a script to be shared with other developers, and the scripting layer, which is the appropriate layer for everyday purposes and is considered a lighter scripting interface to simplify common tasks and for a quick and easy generation of graphics and plots.

Now let's go into each layer in a little more detail:

Back-end layer has three built-in abstract interface classes: FigureCanvas, which defines and encompasses the area on which the figure is drawn. Renderer, an instance of the renderer class knows how to draw on the figure canvas. And finally, Event, which handles user inputs such as keyboard strokes and mouse clicks.

Artist layer: It is composed of one main object, which is the Artist. The Artist is the object that knows how to take the Renderer and use it to put ink on the canvas. Everything you see on a Matplotlib figure is an Artist instance. The title, the lines, the tick labels, the images, and so on, all correspond to an individual Artist. There are two types of Artist objects. The first type is the primitive type, such as a line, a rectangle, a circle, or text. And the second type is the composite type, such as the figure or the axes. The top-level Matplotlib object that contains and manages all of the elements in a given graphic is the figure Artist, and the most important composite artist is the axes because it is where most of the Matplotlib API plotting methods are defined, including methods to create and manipulate the ticks, the axis lines, the grid or the plot background. Now it is important to note that each composite artist may contain other composite artists as well as primitive artists. So, a figure artist for example would contain an axis artist as well as a rectangle or text artists.

Scripting layer: it was developed for scientists who are not professional Programmers. The artist layer is syntactically heavy as it is meant for developers and not for individuals whose goal is to perform quick exploratory analysis of some data. Matplotlib's scripting layer is

essentially the Matplotlib.pyplot interface, which automates the process of defining a canvas and defining a figure artist instance and connecting them.

Read a CSV and Generate a Line Plot with Matplotlib

A line plot is used to represent quantitative values over a continuous interval or time period. It is generally used to depict trends on how the data has changed over time.

In this sub-section, we will see how to use matplotlib to read a csv file and then generate a plot. We will use jupyter notebook. First, we do a basic example to showcase what a line plot is.

```
In [1]: import matplotlib.pyplot as plt
```

The below line will allow you to view the plot inside the jupyter notebook

```
In [2]: %matplotlib inline
```

```
In [3]: import numpy as np
x = np.linspace(0,5,11)
y = (x+1) **2
```

```
In [4]: x
```

```
Out[4]: array([0. , 0.5, 1. , 1.5, 2. , 2.5, 3. , 3.5, 4. , 4.5, 5. ])
```

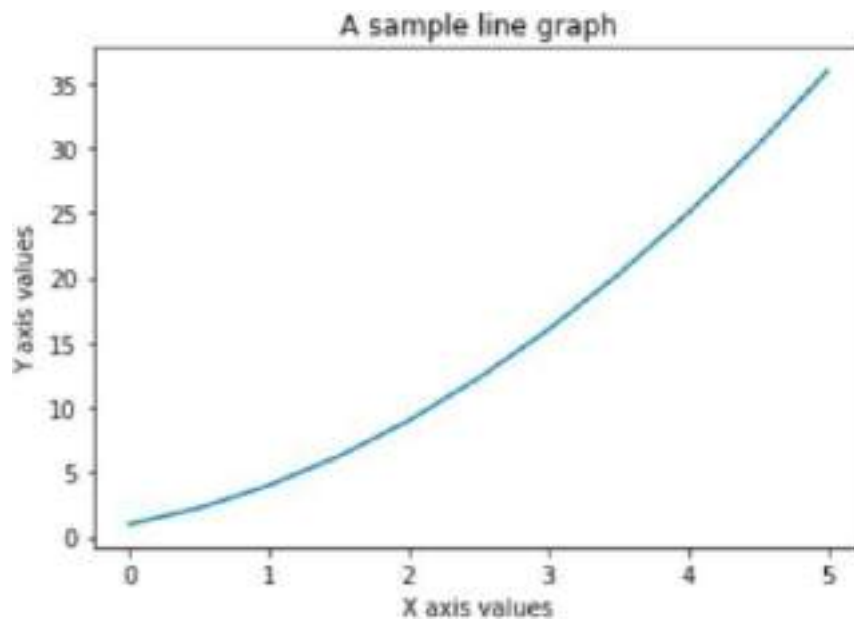
```
In [5]: y
```

```
Out[5]: array([ 1. ,  2.25,  4. ,  6.25,  9. , 12.25, 16. , 20.25, 25. ,
 30.25, 36. ])
```

Now we will plot the graph for these two variables

```
In [9]: plt.plot(x,y)
plt.xlabel('X axis values')
plt.ylabel('Y axis values')
plt.title('A sample line graph')
```

```
Out[9]: Text(0.5, 1.0, 'A sample line graph')
```



Now let us do a small case study using what we just learned now:

Download the dataset from the link:
<https://www.un.org/en/development/desa/population/migration/data/empirical2/migrationflows.asp>

The data set has all the country immigration information. We will use the one for Australia for our case study.

Now let us use this to plot immigration data

```
import numpy as np
import pandas as pd
```

We need to install xrd module that pandas need to read excel files. If you are using anaconda distribution then you can do it by using the command

```
conda install -c anaconda xrd -yes
```

```
df_xus = pd.read_excel('Australia.xlsx', sheet_name='Australia by Residence', skiprows=range(23), skipfooter=2)
print("pandas dataframe contains this data now")
```

pandas dataframe contains this data now

```
df_xus.head()
```

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1988	1989	1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012
6	Emigrants	Both	Afghanistan	035	Asia	5501	Southern Asia	902	Developing regions	0	—	80	120	70	80	120	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
1	Emigrants	Both	Algeria	038	Europe	025	Northern Europe	901	Developed regions	0	—	30	40	30	30	30	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
2	Emigrants	Both	Algeria	033	Africa	112	Northern Africa	902	Developing regions	20	—	20	20	30	40	30	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
3	Emigrants	Both	Armenia	039	Oceania	051	Pacific	902	Developing regions	0	—	0	0	0	0	0	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
4	Emigrants	Both	Austria	038	Europe	025	Southern Europe	901	Developed regions	0	—	0	0	0	0	0	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—

5 rows x 43 columns

```
df_out.head()
```

	Type	Coverage	OdName	AREA	AreaName	REG	RegName	DEV	DevName	1988	1989	1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	2000	2001	2002	2003	2004	2005	2006	2007	2008	2009	2010	2011	2012
441	Immigrants	Both	Wallis and Futuna Islands	909	Oceania	057	Pacific	902	Developing regions	0	—	10	0	0	0	0	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
442	Immigrants	Both	Western Sahara	903	Africa	112	Northern Africa	902	Developing regions	0	—	0	0	0	0	0	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
443	Immigrants	Both	Yemen	038	Asia	022	Western Asia	902	Developing regions	10	—	40	20	10	40	40	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
444	Immigrants	Both	Zambia	903	Africa	113	Eastern Africa	902	Developing regions	150	—	170	240	410	410	400	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—
445	Immigrants	Both	Zimbabwe	903	Africa	113	Eastern Africa	902	Developing regions	030	—	170	150	100	160	220	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—	—

5 rows x 43 columns

Get the list of column names and the list of indices

```
df_aus.columns.values
```

```
array(['Type', 'Coverage', 'OdName', 'AREA', 'AreaName', 'REG', 'RegName',
      'DEV', 'DevName', 1988, 1989, 1990, 1991, 1992, 1993, 1994, 1995, 1996, 1997,
      1998, 1999, 2000, 2001, 2002, 2003, 2004, 2005, 2006, 2007, 2008, 2009,
      2010, 2011, 2012, 2013], dtype=object)
```

```
df_aus.index.values
```

```
array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12,
      13, 14, 15, 16, 17, 18, 19, 20, 21, 22, 23, 24, 25,
      26, 27, 28, 29, 30, 31, 32, 33, 34, 35, 36, 37, 38,
      39, 40, 41, 42, 43, 44, 45, 46, 47, 48, 49, 50, 51,
      52, 53, 54, 55, 56, 57, 58, 59, 60, 61, 62, 63, 64,
      65, 66, 67, 68, 69, 70, 71, 72, 73, 74, 75, 76, 77,
      78, 79, 80, 81, 82, 83, 84, 85, 86, 87, 88, 89, 90,
      91, 92, 93, 94, 95, 96, 97, 98, 99, 100, 101, 102, 103,
      104, 105, 106, 107, 108, 109, 110, 111, 112, 113, 114, 115, 116,
      117, 118, 119, 120, 121, 122, 123, 124, 125, 126, 127, 128, 129,
      130, 131, 132, 133, 134, 135, 136, 137, 138, 139, 140, 141, 142,
      143, 144, 145, 146, 147, 148, 149, 150, 151, 152, 153, 154, 155,
      156, 157, 158, 159, 160, 161, 162, 163, 164, 165, 166, 167, 168,
      169, 170, 171, 172, 173, 174, 175, 176, 177, 178, 179, 180, 181,
      182, 183, 184, 185, 186, 187, 188, 189, 190, 191, 192, 193, 194,
```

Use the `tolist()` method to get index and columns as lists. View the dimensions of dataframe using “.shape” parameter.

After that let us clean the data set to remove few unnecessary columns.

```
df_ase.columns.tolist()
df_ase.index.tolist()

print (type(df_ase.columns.tolist()))
print (type(df_ase.index.tolist()))

class 'list'>
class 'list'>

df_ase.shape
(445, 43)

Let's clean the data set by removing a few unnecessary columns. We can use pandas drop() method as follows:
df_ase.drop(['AREA', 'REG', 'DEV', 'Type', 'Coverage'], axis=1, inplace=True)
df_ase.head(2)
```

	ContName	AreaName	Region	DevName	1980	1981	1982	1983	1984	1985	...	1984	1985	1986	1987	1988	1989	1990	1991	1992	1993
0	Afghanistan	Asia	Southern Asia	Developing regions	0	0	0	0	0	0	...	80	120	70	30	100	-	-	-	-	-
1	Albania	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	30	40	30	30	30	-	-	-	-	-

2 rows x 20 columns

```
df_ase.describe()
```

	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989
count	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000
mean	801.547081	858.704753	942.757198	787.825112	787.483946	705.580538	843.528173	829.717489	1024.753363	1050.822422
std	5282.809981	5381.735219	10381.864208	5320.483640	1853.883474	10354.844888	5116.876347	5487.056488	9054.203458	10180.230194
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000	10.000000
75%	130.000000	140.000000	150.000000	100.000000	100.000000	150.000000	177.500000	190.000000	240.000000	227.500000
max	80810.000000	87680.000000	82340.000000	100510.000000	30360.000000	35448.000000	92400.000000	17770.000000	104770.000000	120040.000000

8 rows x 10 columns

Let us rename the column names so that it makes more sense.

```
df_ase.rename(columns={'ContName':'Country', 'AreaName':'Continent', 'Region':'Region'}, inplace=True)
df_ase.columns

index| 'Country', 'Continent', 'Region', 'DevName', 1980,
      | 1981, 1982, 1983, 1984, 1985,
      | 1986, 1987, 1988, 1989, 1990,
      | 1991, 1992, 1993, 1994, 1995,
      | 1996, 1997, 1998, 1999, 2000,
      | 2001, 2002, 2003, 2004, 2005,
      | 2006, 2007, 2008, 2009, 2010,
      | 2011, 2012, 2013, 'Total'}.
dtype='object')

df_ase.describe()
```

	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989
count	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000	445.000000
mean	811547085	850784753	842737848	787525112	707143896	795560538	883520178	920717189	1024783853	1056022422
std	5232300000	5381755216	5201884280	5328403842	5053363474	5034944999	5156576347	5487018408	6094283458	6516225594
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
max	9004800000	9480200000	9240300000	10851300000	9526000000	9344200000	9245000000	9777000000	15477000000	17004800000

8 rows x 30 columns

```
df_usa.head(3)
```

	Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
0	Afghanistan	Asia	Southern Asia	Developing regions	8	9	0	0	0	0	...	120	79	88	120	-	-	-	-	-	1898
1	Albania	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	40	30	30	30	-	-	-	-	-	1410
2	Algeria	Africa	Northern Africa	Developing regions	20	19	10	10	10	10	...	20	30	40	50	-	-	-	-	-	1170

3 rows x 30 columns

Default index is numerical, but it is more convenient to index based on country names.

```
df_usa.set_index('Country', inplace=True)
df_usa.head(3)
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Country																					
Afghanistan	Asia	Southern Asia	Developing regions	8	9	0	0	0	0	0	...	120	79	88	120	-	-	-	-	-	1898
Albania	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	0	...	40	30	30	30	-	-	-	-	-	1410
Algeria	Africa	Northern Africa	Developing regions	20	19	10	10	10	10	0	...	20	30	40	50	-	-	-	-	-	1170

3 rows x 30 columns

Remove the name of the index.

```
df_usa.index.name = None
```

```
df_usa.head(3)
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Country																					
Afghanistan	Asia	Southern Asia	Developing regions	8	9	0	0	0	0	0	...	120	79	88	120	-	-	-	-	-	1898
Albania	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	0	...	40	30	30	30	-	-	-	-	-	1410
Algeria	Africa	Northern Africa	Developing regions	20	19	10	10	10	10	0	...	20	30	40	50	-	-	-	-	-	1170

3 rows x 30 columns

Let us now test it by pulling the data for Bangladesh.


```
print(df_aus[df_aus.index == 'Bangladesh'].T.squeeze())
```

	Bangladesh	Bangladesh
Continent	Asia	Asia
Region	Southern Asia	Southern Asia
DevName	Developing regions	Developing regions
1980	220	170
1981	140	180
1982	170	220
1983	180	140
1984	160	170
1985	140	220
1986	140	190
1987	140	150
1988	110	270
1989	180	270
1990	130	400
1991	230	600
1992	200	690
1993	190	360
1994	180	690
1995	190	1050
1996	210	900

Column names as numbers could be confusing. For example: year 1985 could be misunderstood as 1985th column. To avoid ambiguity, let us convert column names to strings and then use that to call full range of years.

```
: df_aus.columns = list(map(str, df_aus.columns))
```

```
: years= list(map(str, range(1980,2014)))
years
```

```
: ['1980',
'1981',
'1982',
'1983',
'1984',
'1985',
'1986',
'1987',
'1988',
'1989',
'1990',
'1991',
'1992',
'1993',
'1994',
'1995',
'1996',
'1997',
```

We can also pass multiple criteria in the same line.

```
df_aus[(df_aus['Continent'] == 'Asia') & (df_aus['Region'] == 'Southern Asia')]
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013
Alghanistan	Asia	Southern Asia	Developing regions	0	0	0	0	0	0	0	...	120	70	80	120
Bangladesh	Asia	Southern Asia	Developing regions	220	140	170	180	160	140	140	...	240	400	500	530
Bhutan	Asia	Southern Asia	Developing regions	10	10	10	0	8	0	10	...	50	80	70	100
India	Asia	Southern Asia	Developing regions	740	660	750	580	770	740	760	...	4110	4530	4100	5260
Maldives	Asia	Southern Asia	Developing regions	40	30	30	30	29	50	60	...	100	300	380	350
Maldives	Asia	Southern Asia	Developing regions	10	10	10	0	18	20	10	...	80	110	70	70
Nepal	Asia	Southern Asia	Developing regions	40	70	50	40	80	100	120	...	220	180	180	280
Sumatra	...	Southern	Developing

Let us review the changes we have made to our dataframes.

```
print('data dimension:', df_aus.shape)
print(df_aus.columns)
df_aus.head(5)
```

data dimension: (446, 39)
Index: ['Continent', 'Region', 'DevName', '1980', '1981', '1982', '1983', '1984', '1985', '1986', '1987', '1988', '1989', '1990', '1991', '1992', '1993', '1994', '1995', '1996', '1997', '1998', '1999', '2000', '2001', '2002', '2003', '2004', '2005', '2006', '2007', '2008', '2009', '2010', '2011', '2012', '2013', 'Total'], dtype='object'

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Alghanistan	Asia	Southern Asia	Developing regions	0	0	0	0	0	0	0	...	120	70	80	120	1690
Albania	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	0	...	40	30	10	30	1410

2 rows x 35 columns

Case Study – let us now study the trend of number of immigrants from Bangladesh to Australia.

Case Study - Do a plot of the immigrants from bangladesh

```
bangladesh = df_aus.loc[ 'Bangladesh', years]
bangladesh.head()
```

	1980	1981	1982	1983	1984	1985	1986	1987	1988	1989	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Bangladesh	220	140	170	180	150	140	148	140	110	130	...	480	540	481	500	530
Bangladesh	170	180	220	140	170	220	158	170	270	270	...	2970	1880	2571	1080	3160

2 rows x 34 columns

Since there are two rows of data, let us sum the values of each column and take first 20 years (to eliminate other years for which no values are present).


```
pd_bang=bangladesh.sum(axis=0,skipna=True).head(20)
```

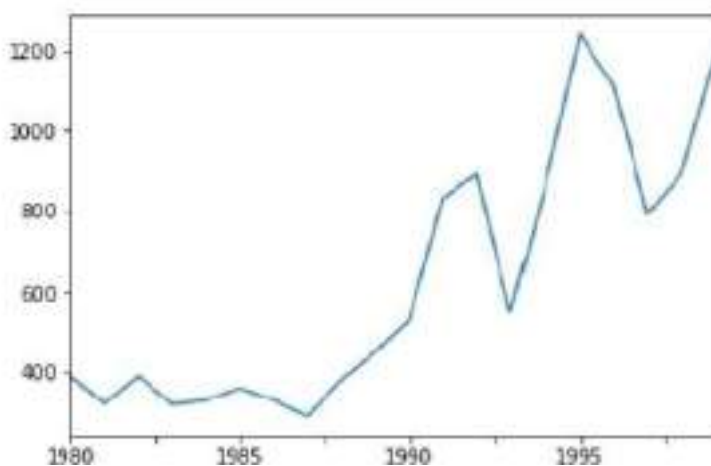
```
pd_bang
```

1980	390
1981	320
1982	390
1983	320
1984	330
1985	360
1986	330
1987	290
1988	380
1989	450
1990	530
1991	830
1992	890
1993	550
1994	870
1995	1240
1996	1110
1997	790
1998	890

Next, we can plot by using the plot function. Automatically the x-axis is plotted with the index values and y-axis with column values

```
pd_bang.plot()
```

```
<matplotlib.axes._subplots.AxesSubplot at 0xc8485c0>
```



Basic Plots using Matplotlib

Area Plot

In the previous module we used line plot to see immigration from Bangladesh to Australia. Now let us try different types of basic plotting using matplotlib.

Area plot

Now let us use area plots to see to visualize cumulative immigration from top 5 countries to Canada. We will use the same process to clean data that we used in the previous section.

```
import numpy as np # useful for many scientific computing in Python
import pandas as pd # primary data structure library
```

URL - https://s3-api.us-geo.objectstorage.softlayer.net/cf-coursesdata/CognitiveClass/DV0101EN/labs/Data_Files/Canada.xlsx

```
df_can = pd.read_excel('https://s3-api.us-geo.objectstorage.softlayer.net/cf-coursesdata/CognitiveClass/DV0101EN/labs/Data_F
                        sheet_name='Canada by Citizenship',
                        skiprows=range(20),
                        skipfooter=1
                        )

print('Data downloaded and read into a dataframe!')
```

Data downloaded and read into a dataframe!

```
df_can.head()
```

	Type	Coverage	OffName	AREA	AreaName	REG	Region	DEV	DevName	1986	1990	1994	1998	2002	2006	2008	2009	2010	2011	2012
0	Immigrants	Foreigners	Algerian	935	Asia	0501	Southern Asia	902	Developing regions	86	—	2979	3406	3009	2952	2111	1746	1758	2203	2636
1	Immigrants	Foreigners	Albania	936	Europe	025	Southern Europe	901	Developed regions	1	—	5480	1223	888	702	560	716	565	539	620
2	Immigrants	Foreigners	Algeria	933	Africa	012	Northern Africa	902	Developing regions	86	—	2616	3026	4887	3023	4885	5383	4752	4325	3774
3	Immigrants	Foreigners	American Samoa	939	Oceania	951	Polynesia	902	Developing regions	6	—	0	0	1	0	0	0	0	0	0
4	Immigrants	Foreigners	Armenia	938	Europe	025	Southern Europe	901	Developed regions	6	—	0	0	1	1	0	0	0	0	1

5 rows x 21 columns

Let's find out how many entries there are in our dataset.

```
# print the dimensions of the dataframe
print(df_can.shape)
```

(155, 21)

Now clean up data using the same process as the one in the previous section :

```
df_can.drop(['AREA', 'POP', 'PCP', 'Type', 'Coverage'], axis=1, inplace=True)
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

	Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
0	Afghanistan	Asia	Southern Asia	Developing regions	18	39	39	47	71	340	...	2978	3420	3009	2852	2111	1746	1756	2203	2535	2094
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	855	702	568	718	581	538	620	683
2	Algeria	Africa	Northern Africa	Developing regions	88	97	71	58	63	48	...	3618	3029	4807	3823	4000	5383	4702	4325	3774	4331
3	American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
4	Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	1

5 rows x 36 columns

```
df_can.rename(columns={'Country':'Country', 'AreaName':'Continent', 'Region':'Region'}, inplace=True)
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

	Country	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
0	Afghanistan	Asia	Southern Asia	Developing regions	18	39	39	47	71	340	...	2978	3420	3009	2852	2111	1746	1756	2203	2535	2094
1	Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	...	1450	1223	855	702	568	718	581	538	620	683
2	Algeria	Africa	Northern Africa	Developing regions	88	97	71	58	63	48	...	3618	3029	4807	3823	4000	5383	4702	4325	3774	4331
3	American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
4	Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	...	0	0	1	1	0	0	0	0	0	1

5 rows x 36 columns

```
# let's examine the types of the column labels
all(isinstance(column, str) for column in df_can.columns)
```

False

```
df_can.columns = list(map(str, df_can.columns))

# let's check the column labels types now
all(isinstance(column, str) for column in df_can.columns)
```

True

```
df_can.set_index('country', inplace=True)
```

```
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

	Continent	Region	Development	1980	1981	1982	1983	1984	1985	1986	...	2004	2005	2006	2007	2008	2009	2010	2011	2012	2013
Country																					
Alghanistan	Asia	Southern Asia	Developing regions	18	39	39	47	71	140	496	...	2978	3455	3869	2052	2111	1746	1758	2203	2635	2904
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	1459	1223	856	732	560	718	561	538	620	503
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	69	...	3618	3620	4807	3623	4605	5393	4752	4325	3774	4331
American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	0	...	0	0	1	0	0	0	0	0	0	0
Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	2	...	0	0	1	1	0	0	0	0	1	1

5 rows × 37 columns

```
df_can['Total'] = df_can.sum(axis=1)
```

```
# let's view the first five elements and see how the dataframe was changed
df_can.head()
```

	Continent	Region	Development	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Country																					
Alghanistan	Asia	Southern Asia	Developing regions	18	39	39	47	71	140	496	...	3435	3869	2052	2111	1746	1758	2203	2635	2904	58610
Albania	Europe	Southern Europe	Developed regions	1	0	0	0	0	0	1	...	1223	856	732	560	718	561	538	620	503	15689
Algeria	Africa	Northern Africa	Developing regions	80	67	71	69	63	44	69	...	3625	4807	3623	4605	5393	4752	4325	3774	4331	80410
American Samoa	Oceania	Polynesia	Developing regions	0	1	0	0	0	0	0	...	0	1	0	0	0	0	0	0	0	8
Andorra	Europe	Southern Europe	Developed regions	0	0	0	0	0	0	2	...	0	1	1	0	0	0	0	1	1	15

5 rows × 38 columns

```
#
```

```
print ('data dimensions:', df_can.shape)
```

```
data dimensions: (195, 38)
```

```
# finally, let's create a list of years from 1980 - 2013
# this will come in handy when we start plotting the data
years = list(map(str, range(1980, 2014)))
```

```
years
```

```
['1980',
 '1981',
 '1982',
 '1983',
 '1984',
 '1985',
 '1986',
 '1987',
 '1988',
```

```
# use the inline backend to generate the plots within the browser
%matplotlib inline

import matplotlib as mpl
import matplotlib.pyplot as plt

mpl.style.use('ggplot') # optional: for ggplot-like style

# check for latest version of Matplotlib
print ('Matplotlib version: ', mpl.__version__) # >= 2.0.0

Matplotlib version: 3.0.3
```

```
df_can.sort_values(['Total'], ascending=False, axis=0, inplace=True)

# get the top 5 entries
df_top5 = df_can.head()

# transpose the dataframe
df_top5 = df_top5[years].transpose()

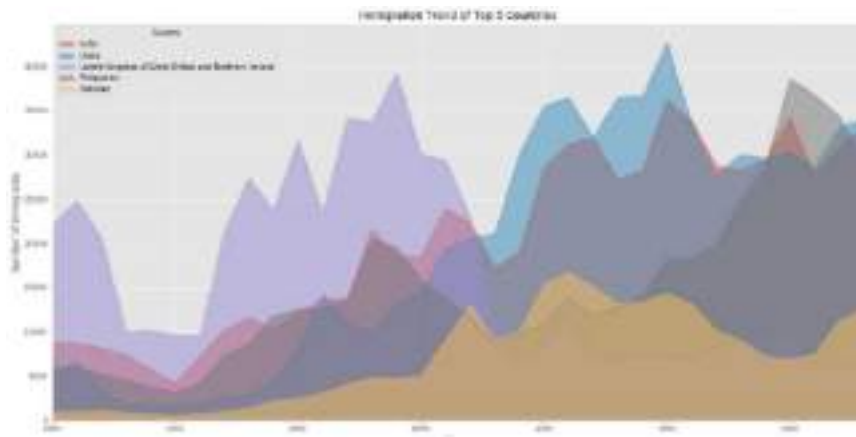
df_top5.head()
```

Country	India	China	United Kingdom of Great Britain and Northern Ireland	Philippines	Pakistan
1980	8880	5123	22045	6051	978
1981	8670	6682	24796	5921	972
1982	8147	3308	20620	5249	1201
1983	7338	1863	10015	4562	900
1984	5704	1527	10170	3801	668

```
df_top5.index = df_top5.index.map(int) # let's change the index values of df_top5 to type integer for plotting
df_top5.plot(kind='area',
             stacked=False,
             figsize=(10, 10), # pass a tuple (x, y) axis
             )

plt.title('Immigration trend of top 5 countries')
plt.ylabel('Number of immigrants')
plt.xlabel('Year')

plt.show()
```

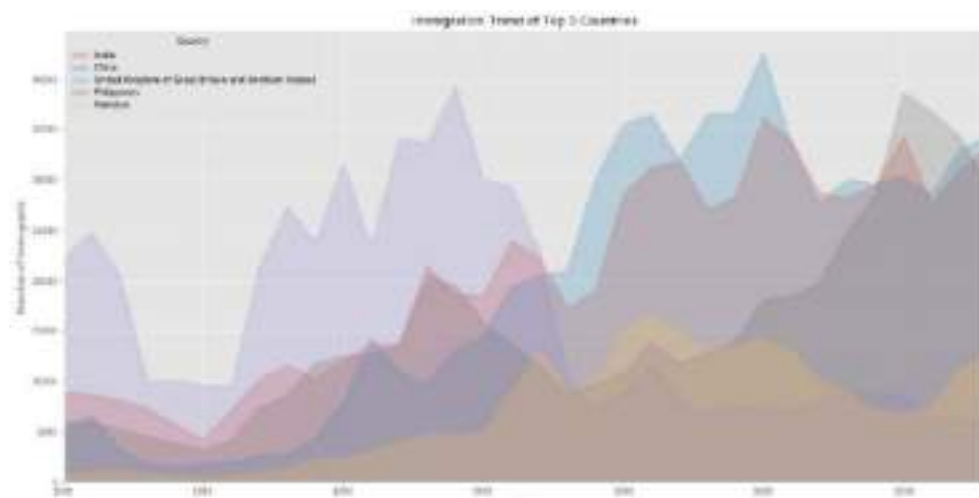


The unstacked plot has a default transparency (alpha value) at 0.5. We can modify this value by passing in the alpha parameter.

```
df_top5.plot(kind='area',
             alpha=0.25, # 0-1, default value is 0.5
             stacked=False,
             figsize=(20, 10),
             )

plt.title('Immigration Trend of Top 5 Countries')
plt.xlabel('Number of Immigrants')
plt.ylabel('Years')

plt.show()
```



Bar Chart

A bar plot is a way of representing data where the length of the bars represents the magnitude/size of the feature/variable. Bar graphs usually represent numerical and categorical variables grouped in intervals.

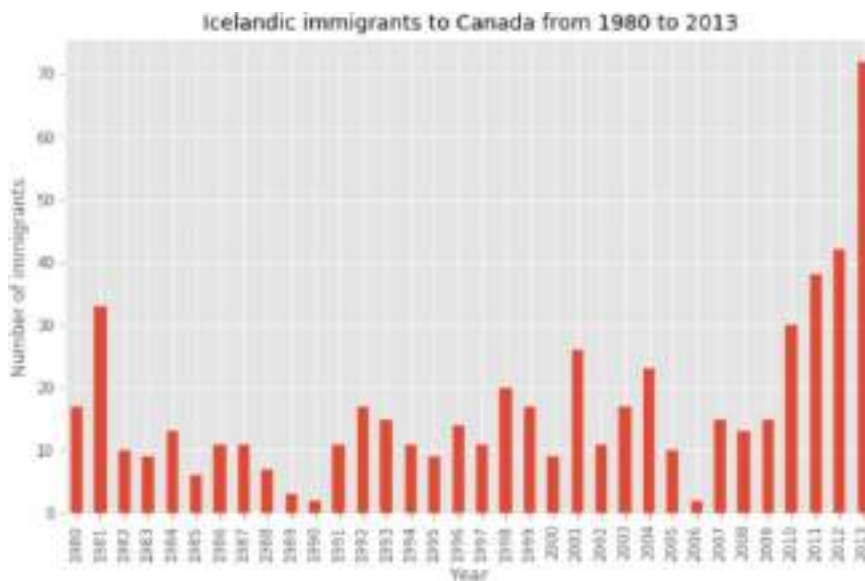
Let's compare the number of Icelandic immigrants (country = 'Iceland') to Canada from year 1980 to 2013.


```
# step 1: get the data
df_iceland = df_can.loc['Iceland', years]
df_iceland.head()
```

```
1980    17
1981    33
1982    10
1983     9
1984    13
Name: Iceland, dtype: object
```

```
# step 2: plot data
df_iceland.plot(kind='bar', figsize=(10, 6))

plt.xlabel('Year') # add to x-label to the plot
plt.ylabel('Number of immigrants') # add y-label to the plot
plt.title('Icelandic immigrants to Canada from 1980 to 2013') # add title to the plot
plt.show()
```



Histogram

How could you visualize the answer to the following question ?

What is the frequency distribution of the number (population) of new immigrants from the various countries to Canada in 2013 ?

To answer this one would need to plot a histogram - it partitions the x-axis into bins, assigns each data point in our dataset to a bin, and then counts the number of data points that have been assigned to each bin. So, the y-axis is the frequency or the number of data points in each

bin. Note that we can change the bin size and usually one needs to tweak it so that the distribution is displayed nicely.

```
# let's quickly view the 2013 data
df_can['2013'].head()
```

```
Country
India                                     33087
China                                    34129
United Kingdom of Great Britain and Northern Ireland    5827
Philippines                             29544
Pakistan                                12603
Name: 2013, dtype: int64
```

```
# np.histogram returns 2 values
count, bin_edges = np.histogram(df_can['2013'])

print(count) # frequency count
print(bin_edges) # bin ranges, default = 10 bins
```

```
[178  11   1   2   0   0   0   0   1   2]
[    0.   3412.9  6825.8 10238.7 13651.6 17064.5 20477.4 23890.3 27303.2
 30716.1 34129. ]
```

By default, the histogram method breaks up the dataset into 10 bins. The figure below summarizes the bin ranges and the frequency distribution of immigration in 2013. We can see that in 2013:

178 Countries contributed between 0 to 3412.9 immigrants

11 Countries contributed between 3412.9 to 6825.8 immigrants

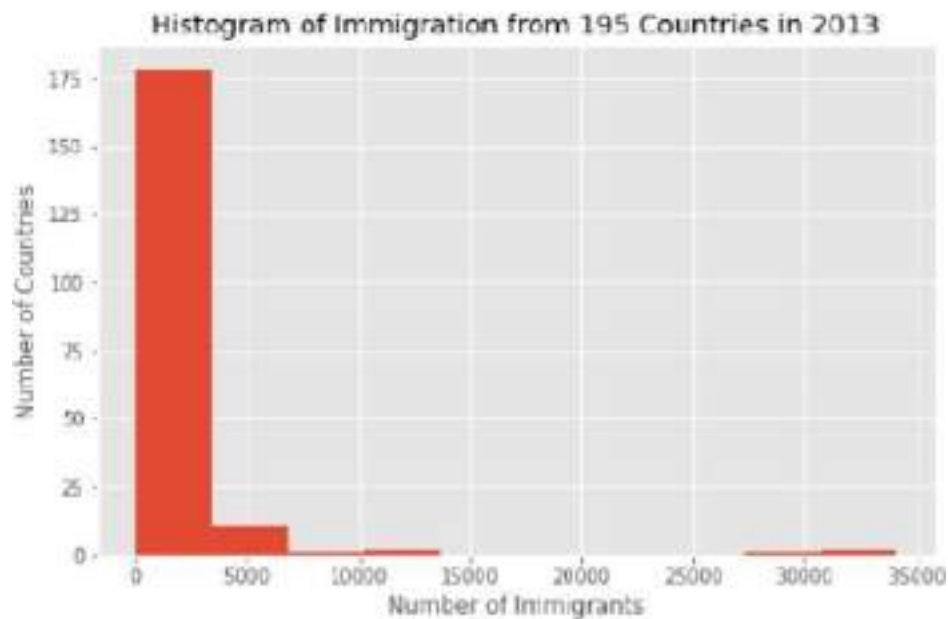
1 Country contributed between 6285.8 to 10238.7 immigrants, and so on.

	Bin 1	Bin 2	Bin 3	Bin 4	Bin 5	Bin 6	Bin 7	Bin 8	Bin 9	Bin 10
Range	0	3412.9	6825.8	10238.7	13651.6	17064.5	20477.4	23890.3	27303.2	30716.1
	to	to	to	to	to	to	to	to	to	to
	3412.9	6825.8	10238.7	13651.6	17064.5	20477.4	23890.3	27303.2	30716.1	34129.
Frequency	178	11	1	2	0	0	0	0	1	2

```
df_can['2013'].plot(kind='hist', figsize=(8, 3))

plt.title('Histogram of Immigration from 199 Countries in 2013') # add a title to the histogram
plt.ylabel('Number of Countries') # add y-label
plt.xlabel('Number of Immigrants') # add x-label

plt.show()
```



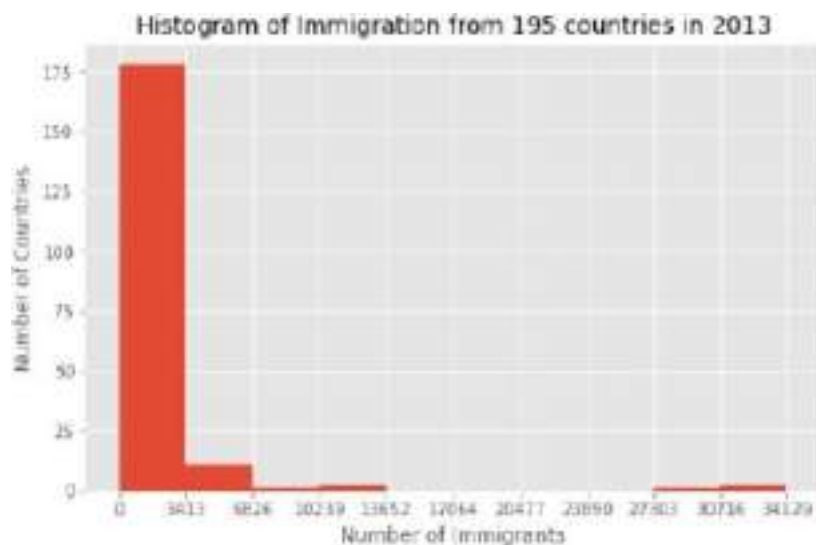
In the above plot, the x-axis represents the population range of immigrants in intervals of 3412.9. The y-axis represents the number of countries that contributed to the population.

Notice that the x-axis labels do not match with the bin size. This can be fixed by passing in a `xticks` keyword that contains the list of the bin sizes, as follows:

```
# 'bin_edges' is a list of bin intervals
count, bin_edges = np.histogram(df_cen['2013'])

df_cen['2013'].plot(kind='hist', figsize=(8, 5), xticks=bin_edges)

plt.title('Histogram of Immigration from 195 countries in 2013') # add a title to the histogram
plt.ylabel('Number of Countries') # add y-label
plt.xlabel('Number of Immigrants') # add x-label
plt.show()
```



Specialized Visualization Tools using Matplotlib

Pie Charts

A pie chart is a circular graphic that displays numeric proportions by dividing a circle (or pie) into proportional slices. You are most likely already familiar with pie charts as it is widely used in business and media. We can create pie charts in Matplotlib by passing in the `kind=pie` keyword.

Let's use a pie chart to explore the proportion (percentage) of new immigrants grouped by continents for the entire time period from 1980 to 2013. We can continue to use the same dataframe further.

```
# group countries by continents and apply size() function
df_continents = df_con.groupby('Continents', axis=0).size()

# note: the output of the groupby method is a 'groupby' object.
# we can not use it further until we Apply a Function (eg. size())
print(type(df_con.groupby('Continents', axis=0)))

df_continents.head()

<class 'pandas.core.groupby.generic.DataFrameGroupby'>
```

	1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	...	2005	2006	2007	2008	2009	2010	2011	2012
Continents																			
Africa	3951	4363	3019	2571	2539	2050	2102	7484	7052	9894	...	27523	29159	29284	29990	34534	60882	35411	36063
Asia	31625	34314	30214	24828	27274	21050	28733	43263	47454	80258	...	180283	149054	133459	130054	141434	163845	146054	152218
Europe	39763	44803	42733	24836	22287	20844	24370	48808	54735	80803	...	36055	33053	33405	34963	35376	33426	26779	29177
Latin America and the Caribbean	13001	11215	10769	16427	13678	10171	21179	23471	21928	25060	...	24747	24676	26711	26547	26867	28819	27836	27173
Northern America	9379	10030	3074	7100	8681	6543	7074	7705	8489	8709	...	8394	9613	3483	10100	2365	8142	7677	7682

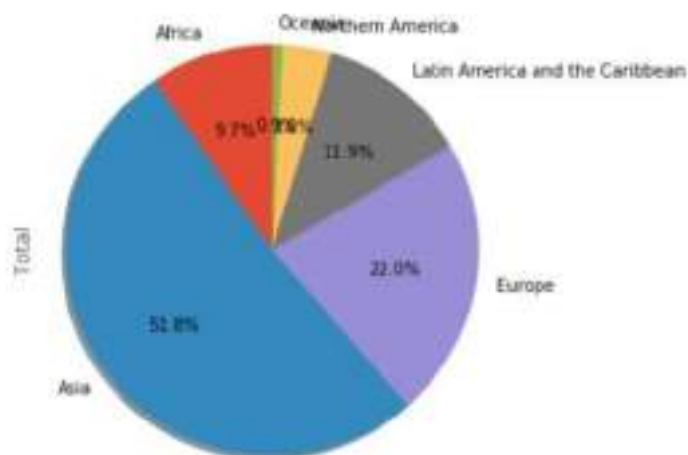
5 rows x 35 columns

```
# autopct create %, start angle represent starting point
df_continents['Total'].plot(kind='pie',
                             figsize=(5, 6),
                             autopct='%1.1f%%', # add in percentages
                             startangle=90,      # start angle 90° (Africa)
                             shadow=True,        # add shadow
                             )

plt.title('Immigration to Canada by Continent [1980 - 2013]')
plt.axis('equal') # Sets the pie chart to look like a circle.

plt.show()
```

Immigration to Canada by Continent [1980 - 2013]



The above visual is not very clear, the numbers and text overlap in some instances. Let's make a few modifications to improve the visuals:

```
colors_list = ['gold', 'yellowgreen', 'lightcoral', 'lightskyblue', 'lightgreen', 'pink']
explode_list = [0.1, 0, 0, 0, 0.1, 0.1] # ratio for each continent with which to offset each wedge.

df_continents['Total'].plot(kind='pie',
                             figsize=(15, 6),
                             autopct='%1.1f%%',
                             startangle=90,
                             shadow=True,
                             labels=None, # turn off labels on pie chart
                             ptdistance=1.1, # the ratio between the center of each pie slice and the start of the text y
                             colors=colors_list, # add custom colors
                             explode=explode_list # 'explode' lowest 3 continents
                            )

# scale the title up by 120 to match ptdistance
plt.title('Immigration to Canada by Continent [1980 - 2013]', y=1.12)

plt.axis('equal')

# add legend
plt.legend(labels=df_continents.index, loc='upper left')

plt.show()
```

Raw code :

```
colors_list = ['gold', 'yellowgreen', 'lightcoral', 'lightskyblue', 'lightgreen', 'pink']
explode_list = [0.1, 0, 0, 0, 0.1, 0.1] # ratio for each continent with which to offset each wedge.
```

```
df_continents['Total'].plot(kind='pie',
                             figsize=(15, 6),
                             autopct='%1.1f%%',
                             startangle=90,
                             shadow=True,
                             labels=None, # turn off labels on pie chart
```

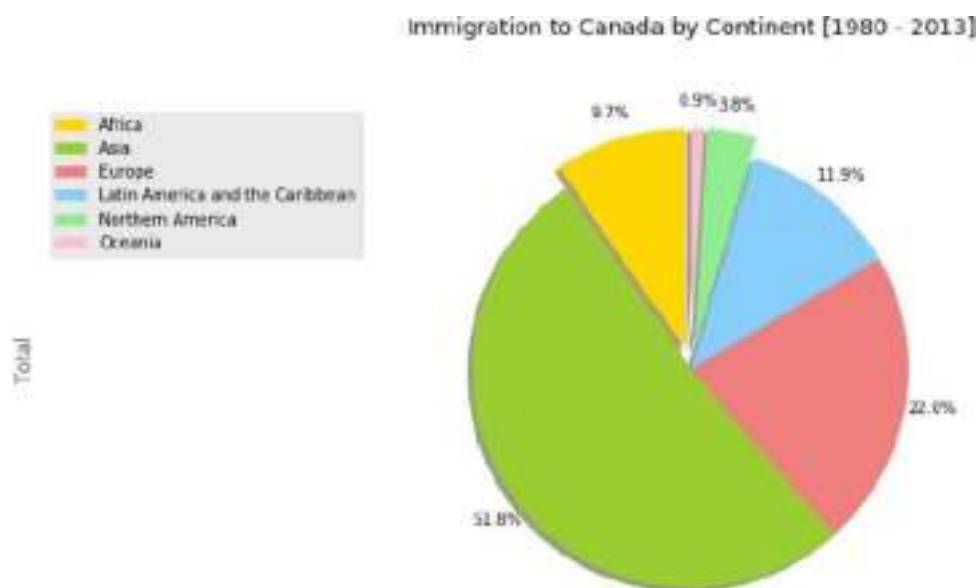
```

pctdistance=1.12, # the ratio between the center of each pie slice and the start of the text
generated by autopct
colors=colors_list, # add custom colors
explode=explode_list # 'explode' lowest 3 continents)

# scale the title up by 12% to match pctdistance
plt.title('Immigration to Canada by Continent [1980 - 2013]', y=1.12)
plt.axis('equal')
# add legend
plt.legend(labels=df_continents.index, loc='upper left')

plt.show()

```



Box Plot

A box plot is a way of statistically representing the distribution of the data through five main dimensions :

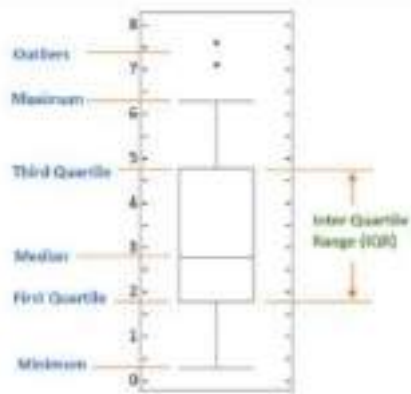
Minimum: Smallest number in the dataset.

First quartile: Middle number between the minimum and the median.

Second quartile (Median): Middle number of the (sorted) dataset.

Third quartile: Middle number between median and maximum.

Maximum: Highest number in the dataset.



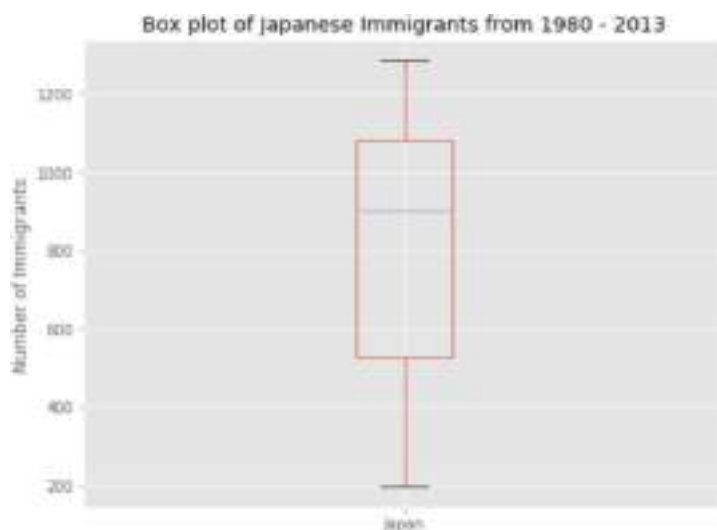
```
# to get a dataframe, place extra square brackets around 'Japan'.
df_japan = df_can.loc[['Japan'], years].transpose()
df_japan.head()
```

Country	Japan
1980	701
1981	756
1982	598
1983	309
1984	246

```
df_japan.plot(kind='box', figsize=(8, 6))

plt.title('Box plot of Japanese Immigrants from 1980 - 2013')
plt.ylabel('Number of Immigrants')

plt.show()
```



We can immediately make a few key observations from the plot above:

The minimum number of immigrants is around 200 (min), maximum number is around 1300 (max), and median number of immigrants is around 900 (median).

25% of the years for period 1980 - 2013 had an annual immigrant count of ~500 or fewer (First quartile).

75% of the years for period 1980 - 2013 had an annual immigrant count of ~1100 or fewer (Third quartile).

We can view the actual numbers by calling the describe() method on the dataframe.

```
df_japan.describe()
```

Country	Japan
count	34.000000
mean	814.911765
std	337.219771
min	198.000000
25%	529.000000
50%	902.000000
75%	1079.000000
max	1284.000000

Scatter Plots

A scatter plot (2D) is a useful method of comparing variables against each other. Scatter plots look similar to line plots in that they both map independent and dependent variables on a 2D graph. While the datapoints are connected by a line in a line plot, they are not connected in a scatter plot. The data in a scatter plot is considered to express a trend. With further analysis using tools like regression, we can mathematically calculate this relationship and use it to predict trends outside the dataset.

Using a scatter plot, let's visualize the trend of total immigration to Canada (all countries combined) for the years 1980 - 2013.

```
# we can use the sum() method to get the total population per year
df_tot = pd.DataFrame(df_can[years].sum(axis=0))

# change the years to type int (useful for regression later on)
df_tot.index = map(int, df_tot.index)

# reset the index to put in back in as a column in the df_tot dataframe
df_tot.reset_index(inplace = True)

# rename columns
df_tot.columns = ['year', 'total']

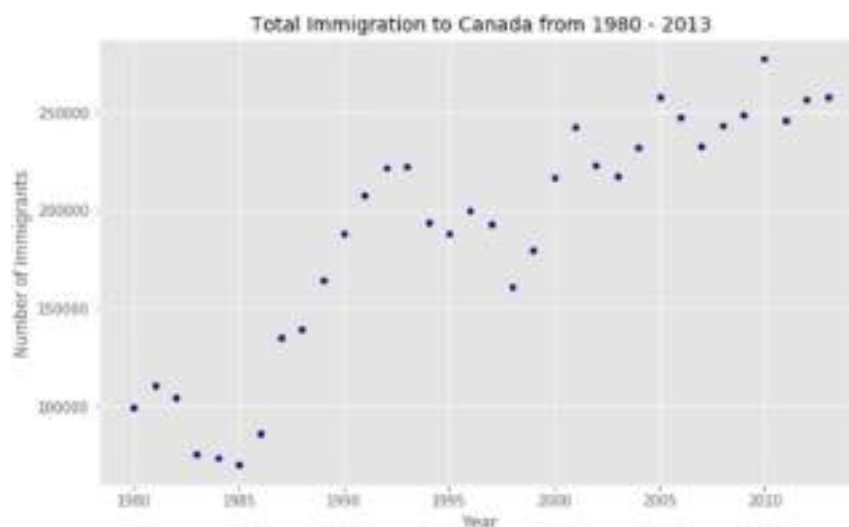
# view the final dataframe
df_tot.head()
```

	year	total
0	1980	99137
1	1981	110563
2	1982	104271
3	1983	75550
4	1984	73417

```
df_tot.plot(kind='scatter', x='year', y='total', figsize=(10, 6), color='darkblue')

plt.title('Total Immigration to Canada from 1980 - 2013')
plt.xlabel('Year')
plt.ylabel('Number of Immigrants')

plt.show()
```



So, let's try to plot a linear line of best fit, and use it to predict the number of immigrants in 2015.

Step 1: Get the equation of line of best fit. We will use Numpy's polyfit() method by passing in the following:

x: x-coordinates of the data.

y: y-coordinates of the data.

deg: Degree of fitting polynomial. 1 = linear, 2 = quadratic, and so on.

```
x = df_tot['year']      # year on x-axis
y = df_tot['total']     # total on y-axis
fit = np.polyfit(x, y, deg=1)

fit

array([ 5.56709228e+03, -1.09261952e+07])
```

Plot the regression line on the scatter plot.

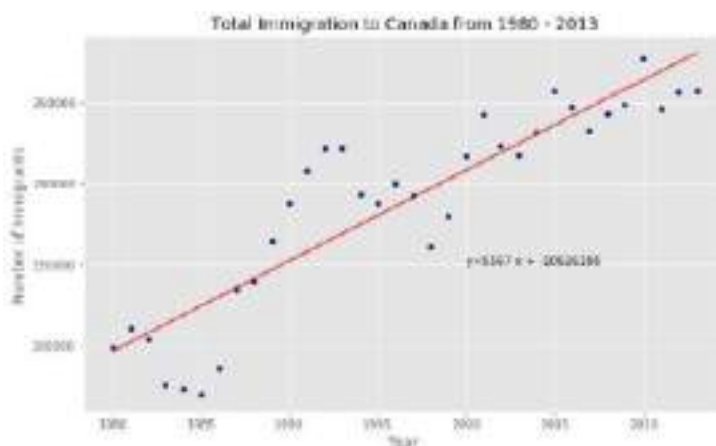
```
df_tot.plot(kind='scatter', x='year', y='total', figsize=(10, 6), color='darkblue')

plt.title('Total Immigration to Canada from 1980 - 2013')
plt.xlabel('Year')
plt.ylabel('Number of Immigrants')

# plot line of best fit
plt.plot(x, fit[0] * x + fit[1], color='red') # recall that x is the Years
plt.annotate('y=(0:.0f) * x + (1:.0f)'.format(fit[0], fit[1]), xy=(2000, 150000))

plt.show()

# print out the line of best fit
'No. Immigrants = (0:.0f) * Year + (1:.0f)'.format(fit[0], fit[1])
```



'No. Immigrants = 5567 * Year + -10926195'

Bubble Plots

A bubble plot is a variation of the scatter plot that displays three dimensions of data (x, y, z). The datapoints are replaced with bubbles, and the size of the bubble is determined by the third variable 'z', also known as the weight. In matplotlib, we can pass in an array or scalar to the keyword `s` to `plot()`, that contains the weight of each point.

Let us compare Argentina's immigration to that of its neighbor Brazil. Let's do that using a bubble plot of immigration from Brazil and Argentina for the years 1980 - 2013. We will set the weights for the bubble as the normalized value of the population for each year.

```
df_can_t = df_can[years].transpose() # transposed dataframe

# cast the Years (the index) to type int
df_can_t.index = asp(int, df_can_t.index)

# let's label the index. This will automatically be the column name when we reset the index
df_can_t.index.name = 'Year'

# reset index to bring the Year in as a column
df_can_t.reset_index(inplace=True)

# view the changes
df_can_t.head()
```

Country	Year	Albania	Algeria	American Samoa	Andorra	Angola	Antigua and Barbuda	Argentina	Armenia	United States of America	Uruguay	Uzbekistan	Vanuatu	Venezuela (Bolivarian Republic of)
0	1980	16	1	60	0	0	1	0	368	0	9378	128	0	0
1	1981	39	0	67	1	0	3	0	426	0	10030	132	0	0
2	1982	39	0	71	0	0	6	0	626	0	9074	146	0	0
3	1983	47	0	69	0	0	6	0	241	0	7100	135	0	0
4	1984	71	0	61	0	0	4	42	237	0	6661	30	0	0

5 rows × 16 columns

Create the normalized weights

There are several methods of normalizations in statistics, each with its own use. In this case, we will use feature scaling to bring all values into the range [0,1]. The general formula is:

where X is an original value, X' is the normalized value. The formula sets the max value in the dataset to 1, and sets the min value to 0. The rest of the datapoints are scaled to a value between 0-1 accordingly.

```
# normalize Brazil data
norm_brazil = (df_can_t['Brazil'] - df_can_t['Brazil'].min()) / (df_can_t['Brazil'].max() - df_can_t['Brazil'].min())

# normalize Argentina data
norm_argentina = (df_can_t['Argentina'] - df_can_t['Argentina'].min()) / (df_can_t['Argentina'].max() - df_can_t['Argentina'].min())
```

Raw Code :

```
# normalize Brazil data
norm_brazil = (df_can_t['Brazil'] - df_can_t['Brazil'].min()) / (df_can_t['Brazil'].max() - df_can_t['Brazil'].min())

# normalize Argentina data
norm_argentina = (df_can_t['Argentina'] - df_can_t['Argentina'].min()) / (df_can_t['Argentina'].max() - df_can_t['Argentina'].min())
```

```
# Brazil
ax0 = df_can_t.plot(kind='scatter',
                    x='Year',
                    y='Brazil',
                    figsize=(14, 8),
                    alpha=0.5, # transparency
                    color='green',
                    s=norm_brazil * 2000 + 10, # pass in weights
                    xlim=(1975, 2015)
                    )

# Argentina
ax1 = df_can_t.plot(kind='scatter',
                    x='Year',
                    y='Argentina',
                    alpha=0.5,
                    color="blue",
                    s=norm_argentina * 2000 + 10,
                    ax = ax0
                    )

ax0.set_ylabel('Number of Immigrants')
ax0.set_title('Immigration from Brazil and Argentina from 1980 - 2013')
ax0.legend(['Brazil', 'Argentina'], loc='upper left', fontsize='x-large')
```

Raw Code :

```
# Brazil
ax0 = df_can_t.plot(kind='scatter',
                    x='Year',
                    y='Brazil',
                    figsize=(14, 8),
                    alpha=0.5, # transparency
```



```

color='green',
s=norm_brazil * 2000 + 10, # pass in weights
xlim=(1975, 2015)
)

# Argentina
ax1 = df_can_t.plot(kind='scatter',
x='Year',
y='Argentina',

                                alpha=0.5,
                                color="blue",
s=norm_argentina * 2000 + 10,
ax = ax0
)

ax0.set_ylabel('Number of Immigrants')
ax0.set_title('Immigration from Brazil and Argentina from 1980 - 2013')
ax0.legend(['Brazil', 'Argentina'], loc='upper left', fontsize='x-large')

```



The size of the bubble corresponds to the magnitude of immigrating population for that year, compared to the 1980 - 2013 data. The larger the bubble, the more immigrants in that year.

Waffle Chart

A waffle chart is an interesting visualization that is normally created to display progress toward goals. It is commonly an effective option when you are trying to add interesting visualization features to a visual that consists mainly of cells, such as an Excel dashboard.

```
# let's create a new dataframe for these three countries
df_dsn = df_can.loc[['Denmark', 'Norway', 'Sweden'], :]

# let's take a look at our dataframe
df_dsn
```

	Continent	Region	DevName	1980	1981	1982	1983	1984	1985	1986	...	2005	2006	2007	2008	2009	2010	2011	2012	2013	Total
Country																					
Denmark	Europe	Northern Europe	Developed regions	272	283	299	306	33	73	93	...	62	101	97	108	81	92	93	94	81	3901
Norway	Europe	Northern Europe	Developed regions	116	77	106	51	31	54	56	...	57	53	73	68	75	48	49	63	59	2127
Sweden	Europe	Northern Europe	Developed regions	281	308	222	176	128	158	187	...	205	131	163	165	107	158	134	140	140	5800

The first step into creating a waffle chart is determining the proportion of each category with respect to the total.

```
# compute the proportion of each category with respect to the total
total_values = sum(df_dsn['Total'])
category_proportions = [(float(value) / total_values) for value in df_dsn['Total']]

# print out proportions
for i, proportion in enumerate(category_proportions):
    print(df_dsn.index.values[i] + ': ' + str(proportion))

Denmark: 0.32255663965602777
Norway: 0.1924094592359848
Sweden: 0.48503390110798744
```

The second step is defining the overall size of the waffle chart.

```
width = 40 # width of chart
height = 10 # height of chart

total_num_tiles = width * height # total number of tiles

print('Total number of tiles is ', total_num_tiles)

Total number of tiles is 400
```

The third step is using the proportion of each category to determine its respective number of tiles

```
# compute the number of tiles for each category
tiles_per_category = [round(proportion * total_num_tiles) for proportion in category_proportions]

# print out number of tiles per category
for i, tiles in enumerate(tiles_per_category):
    print (df_dsn.index.values[i] + ': ' + str(tiles))

Denmark: 129
Norway: 77
Sweden: 194
```

The fourth step is creating a matrix that resembles the waffle chart and populating it.

```
# initialize the waffle chart as an empty matrix
waffle_chart = np.zeros((height, width))

# define indices to loop through waffle chart
category_index = 0
tile_index = 0

# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1

        # if the number of tiles populated for the current category is equal to its corresponding allocated tiles...
        if tile_index > sum(tiles_per_category[0:category_index]):
            # ...proceed to the next category
            category_index += 1

        # set the class value to an integer, which increases with class
        waffle_chart[row, col] = category_index

print ('Waffle chart populated!')

Waffle chart populated!
```

Raw Code :

```
# initialize the waffle chart as an empty matrix
waffle_chart = np.zeros((height, width))
```

```
# define indices to loop through waffle chart
category_index = 0
tile_index = 0
```

```
# populate the waffle chart
for col in range(width):
    for row in range(height):
        tile_index += 1
```

```
# if the number of tiles populated for the current category is equal to its
corresponding allocated tiles...
    if tile_index > sum(tiles_per_category[0:category_index]):
```

```
# ...proceed to the next category
category_index += 1
```

```
# set the class value to an integer, which increases with class
waffle_chart[row, col] = category_index
```

```
print ('Waffle chart populated!')
```

[illegible]

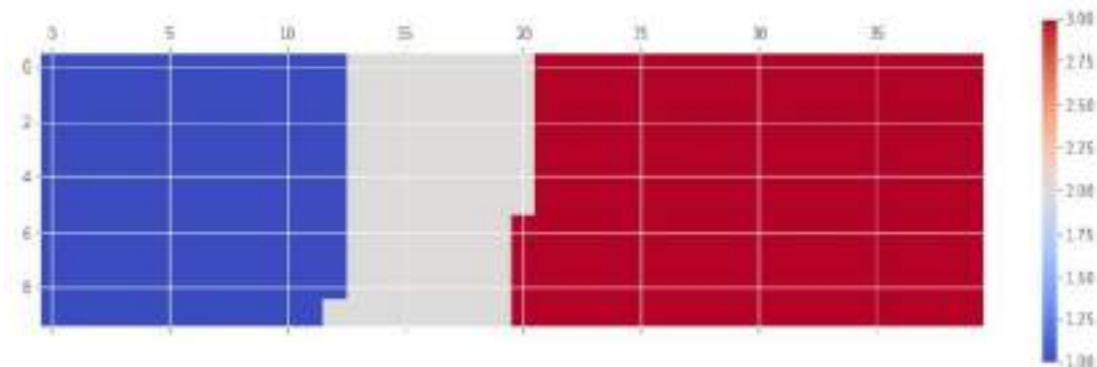
Map the waffle chart matrix into a visual.

```
# instantiate a new figure object
fig = plt.figure()

# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()
```

<matplotlib.colorbar.Colorbar at 0x4076cc0>

<Figure size 432x288 with 0 Axes>



Prettify the chart.

```
# instantiate a new figure object
fig = plt.figure()

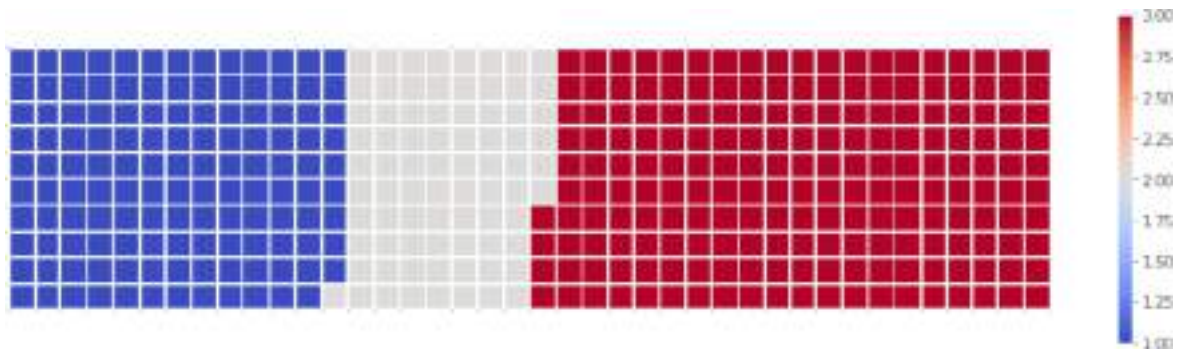
# use matshow to display the waffle chart
colormap = plt.cm.coolwarm
plt.matshow(waffle_chart, cmap=colormap)
plt.colorbar()

# get the axis
ax = plt.gca()

# set minor ticks
ax.set_xticks(np.arange(-.5, (width), 1), minor=True)
ax.set_yticks(np.arange(-.5, (height), 1), minor=True)

# add gridlines based on minor ticks
ax.grid(which='minor', color='w', linestyle='--', linewidth=2)

plt.xticks([])
plt.yticks([])
```



Word Clouds

Word clouds (also known as text clouds or tag clouds) work in a simple way: the more a specific word appears in a source of textual data (such as a speech, blog post, or database), the bigger and bolder it appears in the word cloud.

```
# import package and its set of stopwords
from wordcloud import WordCloud, STOPWORDS

print ('Wordcloud is installed and imported!')
```

Wordcloud is installed and imported!

```
# open the file and read it into a variable alice_novel
alice_novel = open('alice_novel.txt', 'r').read()

print ('File downloaded and saved!')
```

File downloaded and saved!

```
stopwords = set(STOPWORDS)
```

```
# instantiate a word cloud object
alice_wc = WordCloud(
    background_color='white',
    max_words=2000,
    stopwords=stopwords
)

# generate the word cloud
alice_wc.generate(alice_novel)
```



```
# display the word cloud
plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Interesting! So, in the first 2000 words in the novel, the most common words are Alice, said, little, Queen, and so on. Let's resize the cloud so that we can see the less frequent words a little better.

```
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(10) # set height

# display the sliced
plt.imshow(slice_wv, interpolation='bilinear')
plt.axis('off')
plt.show()
```



Much better! However, said isn't really an informative word. So, let's add it to our stop words and re-generate the cloud.

```
stopwords.add('said') # add the words said to stopwords

# re-generate the word cloud
alice_wc.generate(alice_novel)

# display the cloud
fig = plt.figure()
fig.set_figwidth(14) # set width
fig.set_figheight(18) # set height

plt.imshow(alice_wc, interpolation='bilinear')
plt.axis('off')
plt.show()
```

